

Co-occurrence Weighted Conformal Sets for Multi-Label Medical Image Classification

Adrian-Ioan Mosnegutu^a

^a*Faculty of Mathematics and Computer Science, Babes-Bolyai University, Cluj-Napoca, Romania*

Abstract

Conformal prediction provides distribution-free coverage guarantees for uncertainty quantification in deep learning. However, existing multi-label conformal methods either treat labels independently or assume symmetric dependencies, failing to capture the asymmetric relationships between diseases in medical diagnosis. We propose Co-occurrence Weighted Conformal Sets (CWCS), a novel conformal prediction method that models asymmetric label dependencies using empirical disease co-occurrence patterns. CWCS produces prediction sets that respect directional clinical relationships, such as infiltration predicting pneumonia more strongly than the reverse. We benchmark CWCS against standard conformal prediction, nested sets, and tree-based methods on ChestX-ray14 (14 pathology labels, 112,120 images) using a ResNet-50 classifier. Under severe class imbalance (0.2%–17.7% prevalence), all methods achieve valid coverage and similar efficiency. However, CWCS uniquely captures clinically meaningful asymmetric dependencies with measurable directionality, advancing uncertainty quantification for multi-label medical imaging.

Keywords: Conformal prediction, Multi-label classification, Medical imaging, Uncertainty quantification, Chest X-ray, Deep learning, Label dependencies

1. Introduction

Medical image classification represents one of the most critical applications of machine learning, where prediction errors can have severe clinical consequences ranging from delayed treatment to unnecessary interventions. The multi-label nature of medical diagnosis, particularly in chest radiography, presents unique challenges that distinguish it from standard classification tasks. A single chest X-ray may reveal multiple co-occurring pathologies such as cardiomegaly alongside pleural effusion, or infiltration patterns associated with pneumonia. Traditional machine learning approaches typically treat each diagnostic label independently, thereby ignoring the rich structure of dependencies that exists between different medical conditions.

These dependencies are not merely statistical correlations but often reflect underlying pathophysiological mechanisms. For instance, cardiomegaly (an enlarged heart) frequently causes pleural effusion (fluid accumulation around the lungs) through congestive heart failure mechanisms, while infiltration patterns in lung tissue commonly precede or accompany pneumonia. A radiologist naturally considers these relationships when interpreting images, asking questions like “given that I see infiltration, should I look more carefully for signs of pneumonia?” Critically, these relationships are often *asymmetric*: infiltration strongly predicts pneumonia, but the reverse implication is weaker. This directional reasoning suggests that machine learning models would benefit from explicitly modeling such asymmetric dependencies rather than treating each label as an independent prediction problem.

Beyond accurate point predictions, clinical deployment of machine learning systems demands rigorous uncertainty quantification. Conformal prediction (CP) offers an elegant framework for constructing prediction sets with finite-sample coverage guarantees that hold distribution-free, without requiring specific assumptions about the underlying data generating process (Vovk et al., 2005). This property makes conformal prediction particularly attractive for medical applications where distributional assumptions are difficult to verify and model misspecification is common. However, existing conformal prediction methods for multi-label problems present significant limitations. Standard approaches treat labels independently (Angelopoulos et al., 2020), ignoring clinical dependencies entirely. More sophisticated methods model dependencies through symmetric measures (Messoudi et al., 2021) or nested set constructions (Cauchois et al., 2021), but fail to capture the directional nature of many clinical relationships. *No existing method explicitly models asymmetric label dependencies within the conformal prediction framework for multi-label classification.*

Contributions.. This work introduces Co-occurrence Weighted Conformal Sets (CWCS), a novel conformal prediction framework that explicitly leverages asymmetric label dependencies estimated from empirical co-occurrence patterns in training data. Our method constructs prediction sets by incorporating weighted nonconformity scores that account for how strongly the presence of one pathology predicts another, recognizing that this relationship may not be symmetric. We address three key research questions:

1. Can asymmetric dependency modeling produce prediction sets that better align with clinical reasoning while main-

Email address: adrian.mosnegutu@stud.ubbcluj.ro (Adrian-Ioan Mosnegutu)

taining valid coverage guarantees?

2. Under what conditions (particularly severe class imbalance common in medical imaging) do dependency-aware methods provide advantages over simpler baselines?
3. What measurable directionality exists in clinical pathology relationships, and can conformal methods capture it?

Through comprehensive experiments on the ChestX-ray14 dataset encompassing over 112,000 radiographs from more than 30,000 patients, we benchmark CWCS against standard conformal prediction, nested set methods (CDioC), and tree-based approaches modeling symmetric dependencies (CQioC). Our evaluation reveals that under severe class imbalance (pathology prevalence 0.2%–17.7%), all methods achieve similar prediction set efficiency while maintaining valid coverage. However, CWCS uniquely captures clinically meaningful asymmetric dependencies with measurable directionality, producing prediction sets that respect the conditional structure of medical diagnosis. This represents the first systematic investigation of asymmetric dependency modeling in conformal prediction for multi-label medical imaging.

Paper organization. The remainder of this paper is organized as follows. Section 2 reviews related work in uncertainty quantification for medical imaging. Section 3 presents our CWCS method and benchmark approaches. Section 4 describes the experimental setup and ChestX-ray14 dataset. Section 5 presents our CWCS method and benchmark approaches. Section 6 presents comprehensive results and analysis. Section 7 concludes with discussion of limitations and future directions.

2. Related Work

2.1. Conformal Prediction Fundamentals

Conformal prediction, introduced by Vovk et al. (2005), provides a principled framework for uncertainty quantification with finite-sample validity guarantees. The fundamental property states that for a pre-specified miscoverage rate $\alpha \in (0, 1)$, the constructed prediction set contains the true label with probability at least $1 - \alpha$, and this guarantee holds regardless of the underlying data distribution. This distribution-free characteristic makes conformal prediction particularly valuable in medical applications where we cannot rely on specific parametric assumptions.

The split conformal prediction approach (Papadopoulos et al., 2002) divides available data into a proper training set for learning the base predictor and a separate calibration set for computing nonconformity scores. These scores quantify how unusual or “nonconforming” each prediction would be relative to the calibration distribution. By using the empirical quantile of calibration scores, the method constructs prediction sets that provably satisfy the coverage guarantee. Recent theoretical advances have extended these ideas to handle more complex scenarios including conformalized quantile regression (Romano et al., 2019) and uncertainty sets for neural network classifiers (Angelopoulos et al., 2020).

2.2. Multi-Label Conformal Prediction

The extension of conformal prediction to multi-label classification introduces unique challenges compared to standard multi-class problems. The most straightforward approach treats each label independently, constructing separate prediction intervals for each class (Sadinle et al., 2019). While simple and computationally efficient, this independence assumption ignores valuable structural information about label relationships and can lead to statistically inefficient or clinically implausible prediction sets.

Recognizing these limitations, several researchers have proposed dependency-aware approaches. Conditional Conformal Prediction with Increasing Order of Conditioning (CDioC) uses a predetermined ordering of labels to sequentially condition predictions (Stutz et al., 2021), with the Adaptive Prediction Sets (APS) variant (Romano et al., 2020) constructing nested prediction sets based on cumulative probability mass. This approach has shown promise in multi-class problems by adaptively adjusting set sizes based on instance difficulty. However, the extension to multi-label settings with sparse ground truth presents challenges, as we will demonstrate in our experimental results.

For structured prediction problems, graphical models provide a natural framework for representing dependencies. The Tree-based Conformal Quantile of Inclusion Order (CQioC) approach (Fontana et al., 2023) employs the Chow-Liu algorithm to learn a maximum spanning tree over labels, where edges represent pairwise dependencies weighted by mutual information. This tree structure enables efficient inference while respecting learned dependencies. However, because mutual information is inherently symmetric, tree-based methods cannot capture directional relationships where the conditional probability of label B given label A differs substantially from the conditional probability of label A given label B.

2.3. Medical Image Classification and Uncertainty

Multi-label classification of medical images has been extensively studied, particularly for chest radiography. The ChestX-ray14 dataset (Wang et al., 2017), consisting of 112,120 frontal-view chest X-rays labeled with 14 thoracic pathology classes through automated text mining of radiology reports, has become a standard benchmark. Subsequent datasets like CheXpert (Irvin et al., 2019) have introduced uncertainty labels to better capture the ambiguity inherent in radiological interpretation, while recent work on GLoRIA (Huang et al., 2021) has demonstrated the benefits of multi-modal learning that leverages both images and associated radiology reports.

Despite impressive advances in classification accuracy through modern architectures like vision transformers (Dosovitskiy et al., 2021) and contrastive learning approaches (Chen et al., 2020), uncertainty quantification in medical imaging has received less attention. Bayesian approaches using Monte Carlo dropout (Gal and Ghahramani, 2016) and deep ensembles (Lakshminarayanan et al., 2017) can provide uncertainty estimates but lack theoretical guarantees. Recent work has begun exploring conformal prediction for medical applications including segmentation tasks (Angelopoulos et al., 2022) and fairness

considerations (Lu et al., 2022), but the specific challenge of modeling asymmetric label dependencies in multi-label medical classification remains largely unaddressed.

The class imbalance problem in medical datasets presents particular challenges for both classification and uncertainty quantification. In ChestX-ray14, pathology prevalences span nearly two orders of magnitude, from Hernia appearing in only 0.2% of images to Infiltration present in 17.7%. This severe imbalance means that standard classification metrics can be misleading, and achieving well-calibrated probability estimates for rare classes becomes extremely difficult. Understanding how conformal prediction methods behave under such extreme imbalance is crucial for practical deployment.

3. Background and Problem Formulation

3.1. Notation and Setup

Let $\mathcal{X}$ denote the space of input images (chest X-rays) and $\mathcal{Y} = \{0, 1\}^L$ denote the space of multi-label outputs, where L represents the number of possible pathology labels. For the ChestX-ray14 dataset, we have $L = 14$ corresponding to pathologies including Atelectasis, Cardiomegaly, Effusion, Infiltration, Mass, Nodule, Pneumonia, Pneumothorax, Consolidation, Edema, Emphysema, Fibrosis, Pleural Thickening, and Hernia. Each image $X \in \mathcal{X}$ is associated with a binary label vector $Y \in \mathcal{Y}$ where $Y^{(j)} = 1$ indicates the presence of pathology j and $Y^{(j)} = 0$ indicates its absence.

We observe independent and identically distributed samples $\{(X_i, Y_i)\}_{i=1}^n$ drawn from an unknown joint distribution P_{XY} . Our goal is to construct a prediction set function $C : \mathcal{X} \rightarrow 2^{\{1, \dots, L\}}$ that maps each input image to a subset of labels. This prediction set should satisfy the coverage property that with high probability, all true labels are contained within the set. Formally, we require instance-level coverage defined as:

$$\mathbb{P}_{(X,Y) \sim P_{XY}}(Y^{(j)} = 1 \text{ for all } j \in C(X)) \geq 1 - \alpha \quad (1)$$

for a pre-specified miscoverage rate $\alpha \in (0, 1)$. In our experiments, we target $\alpha = 0.10$, corresponding to 90% coverage. We also consider label-level coverage, defined separately for each label j as the proportion of instances where that label, if present in the ground truth, is included in the prediction set.

3.2. Split Conformal Prediction Framework

The split conformal prediction protocol operates through a carefully designed sequence of steps that separate model training from calibration. We partition the available data into three disjoint sets serving distinct purposes. The training set $\mathcal{D}_{\text{train}}$ of size n_{train} is used to learn the base probabilistic classifier $\hat{\pi} : \mathcal{X} \rightarrow [0, 1]^L$, where $\hat{\pi}^{(j)}(X)$ provides an estimate of the conditional probability $P(Y^{(j)} = 1 | X)$. The calibration set $\mathcal{D}_{\text{cal}} = \{(X_i, Y_i)\}_{i=1}^{n_{\text{cal}}}$ is held out during training and used exclusively for computing nonconformity scores and calibration quantiles. Finally, the test set $\mathcal{D}_{\text{test}}$ is reserved for final evaluation and never used during either training or calibration.

The core of conformal prediction lies in the construction and use of nonconformity scores. For each calibration sample (X_i, Y_i) , we compute a scalar score s_i that measures how unusual or nonconforming the true label Y_i is relative to the predicted probability distribution $\hat{\pi}(X_i)$. The specific definition of this score varies across different conformal methods and represents the key design choice that distinguishes one approach from another. After computing scores for all calibration samples, we determine the calibrated quantile by taking the $(1 - \alpha)$ -quantile of the augmented score distribution that includes an additional infinite score:

$$\hat{q} = \text{Quantile}(\{s_1, \dots, s_{n_{\text{cal}}}, \infty\}, 1 - \alpha) \quad (2)$$

The inclusion of the infinite score, while seemingly technical, is crucial for achieving exact finite-sample coverage rather than asymptotic coverage. For a test instance X_{test} , we construct the prediction set by including all labels whose nonconformity scores fall below or equal to this calibrated threshold.

Under the assumption that the data are exchangeable (which holds trivially when they are independent and identically distributed), this procedure provides the guarantee that the probability of the true label being contained in the prediction set is at least $1 - \alpha$. Importantly, this guarantee holds for any base classifier $\hat{\pi}$, regardless of whether it is well-calibrated or even accurate. The conformal framework wraps around the base classifier to provide valid uncertainty quantification even when the underlying model is imperfect.

3.3. Modeling Label Dependencies

In multi-label medical classification, labels are rarely independent. We quantify these relationships through the empirical co-occurrence matrix $\mathbf{M} \in [0, 1]^{L \times L}$, where each entry represents a conditional probability:

$$M_{jk} = \mathbb{P}(Y^{(k)} = 1 | Y^{(j)} = 1) = \frac{\sum_{i=1}^n Y_i^{(j)} \cdot Y_i^{(k)}}{\sum_{i=1}^n Y_i^{(j)}} \quad (3)$$

This matrix is directional and generally asymmetric. The entry M_{jk} answers the question: “given that pathology j is present, what is the probability that pathology k is also present?” This is distinct from M_{kj} , which answers the reverse question. The asymmetry between these two quantities provides crucial information about directional relationships.

Definition 1 (Asymmetric Dependency). *For any pair of labels j and k , we define the asymmetry measure as:*

$$\Delta_{jk} = |M_{jk} - M_{kj}| \quad (4)$$

A large value of Δ_{jk} indicates a strong directional relationship where one label is a much better predictor of the other than vice versa.

These asymmetries are particularly pronounced in medical contexts where causal or sequential relationships exist between pathologies. For example, in our analysis of ChestX-ray14, we observe that infiltration strongly predicts pneumonia with

$M_{\text{Inf,Pneu}} = 0.42$, meaning that 42% of images showing infiltration also show pneumonia. However, the reverse relationship is much weaker, with $M_{\text{Pneu,Inf}} = 0.15$, yielding an asymmetry of $\Delta = 0.27$. This pattern reflects clinical reality: infiltration (abnormal density in lung tissue) is a common radiological finding in pneumonia, but infiltration has many non-infectious causes, so the presence of pneumonia provides relatively weak evidence for infiltration when considering all possible etiologies.

4. Experimental Methodology

4.1. Dataset Description and Characteristics

We conduct our experiments on the ChestX-ray14 dataset, which represents one of the largest publicly available collections of annotated chest radiographs. The dataset comprises 112,120 frontal-view X-ray images acquired from 30,805 unique patients at the National Institutes of Health Clinical Center. All images are grayscale with varying resolutions, typically ranging from approximately 1024 by 1024 pixels to 2500 by 2500 pixels. The dataset includes comprehensive metadata for each image including patient age, gender, view position (PA or AP), and the original image dimensions.

The pathology labels were extracted from radiological reports through an automated natural language processing pipeline that searched for mentions of 14 specific thoracic diseases. It is important to note that this automated extraction process introduces some level of label noise, with previous studies estimating the accuracy at approximately 90 to 95 percent. This noise ceiling represents a fundamental limitation on achievable performance, but the large scale of the dataset makes it valuable for developing and evaluating multi-label classification methods despite this imperfection.

The class imbalance in this dataset is severe, which reflects both the natural prevalence of different conditions in the patient population and the challenges of detecting rare pathologies in automated label extraction. Figure 1 illustrates this distribution, showing that common conditions like Infiltration appear in nearly one out of every five images, while rare conditions like Hernia occur in fewer than one out of every 400 images. This 88-fold difference in prevalence creates fundamental challenges for both classifier training and conformal calibration, as rare classes have insufficient samples for precise probability estimation.

Beyond individual label frequencies, the multi-label distribution reveals that most images are either normal or contain only one or two pathologies. Approximately 54% of images in the dataset are labeled as “No Finding,” indicating no significant abnormalities detected. Among pathological cases, the median number of concurrent conditions is one, with only about 6% of images showing three or more simultaneous pathologies. This sparsity has important implications for evaluating conformal prediction methods: when the ground truth contains zero or one label but prediction sets contain nine or ten labels, this represents a substantial mismatch that we must carefully interpret.

4.2. Data Preprocessing and Splitting Strategy

A critical methodological consideration in medical imaging studies is ensuring that the same patient does not appear in multiple data partitions, as this would lead to optimistically biased performance estimates due to learning patient-specific characteristics rather than generalizable disease patterns. We implement strict patient-level splitting by first grouping all 112,120 images by their associated patient identifier, yielding 30,805 unique patients. We then randomly shuffle these patients using a fixed random seed (2024 in our experiments) and assign them to partitions according to the following ratios: 70% for training (21,563 patients corresponding to 78,440 images), 10% for validation (3,080 patients corresponding to 11,244 images), 10% for calibration (3,080 patients corresponding to 10,954 images), and 10% for testing (3,082 patients corresponding to 11,482 images).

The validation set serves a distinct purpose from the calibration set in our experimental protocol. We use the validation set during base classifier training for hyperparameter tuning, early stopping decisions, and monitoring training progress through metrics like loss and F1 score. The validation set is never used for conformal prediction calibration. In contrast, the calibration set is held completely separate during all phases of classifier training and is used exclusively for computing nonconformity scores and calibration quantiles in the conformal prediction framework. This strict separation is essential for maintaining the theoretical validity of conformal coverage guarantees. Finally, the test set is reserved for reporting all final results and is never accessed during any training or calibration procedures.

For image preprocessing, we apply a standardized pipeline to prepare inputs for the ResNet-50 backbone network. All images are resized to 224 by 224 pixels using bilinear interpolation to match the expected input dimensions of the pretrained ImageNet model. Since ResNet-50 expects three-channel RGB inputs but chest X-rays are grayscale, we replicate the single intensity channel across all three color channels. We then normalize pixel intensities using ImageNet statistics with means of 0.485, 0.456, and 0.406 and standard deviations of 0.229, 0.224, and 0.225 for the three channels respectively. During training, we apply data augmentation including random horizontal flips with probability 0.5, random rotations within plus or minus 10 degrees, and random brightness and contrast adjustments. Importantly, we do not apply any augmentation during validation, calibration, or testing to ensure that nonconformity scores reflect the true data distribution.

4.3. Base Classifier Architecture and Training

We employ ResNet-50 (He et al., 2016), a widely-used convolutional neural network architecture, as our base classifier. ResNet-50 consists of 50 layers organized into four main blocks with residual connections that enable training of deep networks by mitigating vanishing gradient problems. We initialize the network with weights pretrained on ImageNet, leveraging transfer learning to benefit from features learned on natural images. The original classification head designed for 1000-way ImageNet classification is replaced with a custom head suitable




Figure 1: Label distribution characteristics of ChestX-ray14. The left panel shows pathology frequencies in the training set, revealing extreme class imbalance spanning two orders of magnitude from Hernia at 0.2% prevalence to Infiltration at 17.7%. The right panel displays the distribution of number of labels per image, demonstrating that the majority of images (54%) are labeled as “No Finding” with zero pathologies, while most pathological cases present with one or two concurrent conditions. This sparsity pattern has important implications for conformal prediction, as methods must maintain coverage when ground truth contains zero or few labels while prediction sets may contain many labels.

for multi-label prediction consisting of a dropout layer with probability 0.5 for regularization followed by a linear layer projecting from 2048 features to 14 output logits. We apply sigmoid activation to produce independent probabilities for each label, which is essential for multi-label classification as it allows multiple labels to be predicted simultaneously without forcing them into a probability distribution that sums to one.

Our training strategy follows a two-phase approach designed to effectively leverage transfer learning. In the first phase, we freeze all parameters of the ResNet-50 backbone and train only the dropout and final linear layer for 30 epochs. This allows the classification head to adapt to our specific task while preserving the general-purpose features learned from ImageNet. We use the Adam optimizer with learning rate of 0.001 and weight decay of 0.00001 for regularization. The loss function is binary cross-entropy computed independently for each label, treating the multi-label problem as 14 separate binary classification tasks. We use a batch size of 32 images and clip gradients to a maximum norm of 1.0 to prevent exploding gradients during training.

In the second phase, we unfreeze layers 3 and 4 of the ResNet-50 backbone while keeping the earlier layers frozen. We fine-tune these layers along with the classification head for an additional 10 epochs using a reduced learning rate of 0.0001 to avoid catastrophic forgetting of the pretrained features. This two-phase strategy balances adaptation to our specific task with preservation of useful pretrained representations. We implement early stopping with patience of 5 epochs based on validation set macro-F1 score, although in our experiments the full 30 plus 10 epochs were typically needed before convergence.

The base classifier achieves a macro-F1 score of 0.32 on the validation set and test set. While this performance may appear modest, it is important to understand this result in context. The macro-F1 metric equally weights all 14 pathology

classes, meaning that poor performance on rare classes like Hernia and Pneumonia substantially reduces the overall score even if the model performs well on common classes. Given that Hernia appears in only 140 images in the training set out of over 78,000 total training images, it is extremely difficult for the model to learn robust features for this class. Additionally, the label noise from automated text extraction provides a fundamental ceiling on achievable performance. Figure 2 shows that the model continues to improve on the training set throughout training while validation performance plateaus, suggesting that further gains would require addressing the class imbalance through techniques like oversampling rare classes or using specialized loss functions, or improving label quality through manual annotation. For our purposes, this moderate-performance classifier serves as a realistic baseline for evaluating conformal prediction methods under challenging practical conditions.

5. Conformal Prediction Methods

5.1. Standard Conformal Prediction with Independent Calibration

The simplest approach to multi-label conformal prediction treats each label independently, computing separate calibration thresholds for each of the 14 pathology classes. For a given label j , the nonconformity score for a calibration sample (X_i, Y_i) is defined as one minus the predicted probability for that label:

$$s_i^{(j)} = 1 - \hat{\pi}^{(j)}(X_i) \quad (5)$$

This score is high when the model assigns low probability to the label, making it “nonconforming” or unusual. Crucially, we only compute and store this score for calibration samples where the label is actually present, that is, where $Y_i^{(j)} = 1$. This ensures that our calibration quantile represents the threshold needed to achieve coverage on positive instances of each label.




Figure 2: Training curves for the ResNet-50 base classifier over 30 epochs. The left panel shows training loss (blue) steadily decreasing from 0.80 to 0.49, while validation loss (orange) decreases initially but plateaus around 0.71 after epoch 20, with slight subsequent degradation indicating mild overfitting. The right panel displays validation macro-F1 score (green) progression, reaching a maximum of approximately 0.28 around epochs 22 through 25. The target F1 of 0.40 (red dashed line) is not achieved, reflecting fundamental challenges from extreme class imbalance where rare pathologies like Hernia have only 140 training examples and label noise from automated text extraction affecting 5 to 10 percent of annotations. This moderate performance (F1 equals 0.32 on test set) necessitates conservative conformal calibration regardless of dependency modeling approach.

For each label j independently, we collect all nonconformity scores from calibration samples where that label is present, forming the set $\mathcal{S}^{(j)} = \{s_i^{(j)} : Y_i^{(j)} = 1, i = 1, \dots, n_{\text{cal}}\}$. We then compute the calibrated threshold for label j as:

$$\hat{q}^{(j)} = \text{Quantile}(\mathcal{S}^{(j)} \cup \{\infty\}, 1 - \alpha) \quad (6)$$

At test time, label j is included in the prediction set if its nonconformity score is at or below its calibrated threshold, equivalently, if its predicted probability exceeds the threshold $\tau^{(j)} = 1 - \hat{q}^{(j)}$. The final prediction set is formed as:

$$C(X_{\text{test}}) = \{j \in \{1, \dots, 14\} : \hat{\pi}^{(j)}(X_{\text{test}}) \geq \tau^{(j)}\} \quad (7)$$




Figure 3: Calibrated probability thresholds for Standard CP across all 14 pathology labels. Most labels require very low thresholds ranging from 0.02 to 0.24, far below the conventional 0.5 decision boundary (shown as red dashed line). These low thresholds reflect the poor calibration of the base classifier under severe class imbalance. For instance, Infiltration requires a threshold of 0.24, meaning any test image where the classifier assigns Infiltration probability exceeding 24% will have Infiltration included in its prediction set. When the classifier frequently assigns multiple labels probabilities in the 0.20 to 0.40 range simultaneously due to uncertainty, this leads to the large prediction sets observed in our experiments averaging nearly 10 labels per image.

Figure 3 reveals that Standard CP requires very low inclusion thresholds for most pathologies, typically between 0.02 and

0.24. This reflects the fundamental challenge of poor probability calibration under class imbalance. To guarantee that 90% of positive instances of each rare class are captured, the method must set thresholds conservatively low. When applied independently across all 14 labels, this results in large prediction sets even though Standard CP achieves the smallest average set size among all methods we evaluate, as it does not boost labels based on co-occurrence with already-included labels.

5.2. Adaptive Prediction Sets (APS)

The Adaptive Prediction Sets approach, originally developed for multi-class classification by Romano et al. (2020), constructs prediction sets by ordering labels according to predicted probability and including labels sequentially until a calibrated threshold is exceeded. We adapt this method to the multi-label setting following the approach of Stutz et al. (2021).

For a calibration sample (X_i, Y_i) , we first sort all labels in descending order of predicted probability, obtaining a permutation σ_i where $\sigma_i(1)$ is the most confident label, $\sigma_i(2)$ is the second most confident, and so on. We then compute the cumulative sum of predicted probabilities in this sorted order. The nonconformity score is defined as the cumulative probability mass that must be included before capturing all true labels:

$$s_i = \sum_{k=1}^{K_i} \hat{\pi}^{(\sigma_i(k))}(X_i) \quad (8)$$

where K_i is the smallest index such that all true labels (those where $Y_i^{(j)} = 1$) are included in the set $\{\sigma_i(1), \dots, \sigma_i(K_i)\}$. Intuitively, this score measures how much probability mass we must accumulate before capturing all true positives. Samples where true labels receive low probabilities will have high nonconformity scores because we must include many labels to eventually capture them all.

After computing scores for all calibration samples, we determine the calibrated quantile $\hat{q} = \text{Quantile}(\{s_1, \dots, s_{n_{\text{cal}}}, \infty\}, 1 - \alpha)$ as in standard conformal prediction. For a test instance, we construct the prediction set by including labels in decreasing probability order until the cumulative probability mass first exceeds $\hat{q}$:

$$C(X_{\text{test}}) = \{\sigma(1), \dots, \sigma(K)\} \quad (9)$$

where $K = \min \{k : \sum_{j=1}^k \hat{\pi}^{(\sigma(j))}(X_{\text{test}}) > \hat{q}\}$. This greedy construction ensures that prediction sets are nested: if label k is excluded, all labels with lower probabilities are also excluded. This nesting property is important for maintaining valid coverage under the conformal framework.

In practice, we find that APS produces substantially larger prediction sets than other methods in our experiments, averaging 11.80 labels per image compared to approximately 10 for other approaches. This behavior stems from how APS handles sparse multi-label ground truth. When the majority of images have zero or one true pathology, but the classifier's probability is spread across multiple labels due to uncertainty, the cumulative mass threshold becomes conservative to ensure coverage. The method effectively says “include labels until you have accumulated enough probability mass to be confident that true labels are captured,” and under poor calibration, this mass accumulation requires including many labels.

5.3. Tree-based Conformal Prediction with Chow-Liu Trees

Tree-based conformal methods model label dependencies through a learned tree structure where nodes represent labels and edges represent dependencies. The Chow-Liu algorithm (Chow and Liu, 1968) finds the maximum spanning tree that best approximates the joint distribution of labels using pairwise mutual information as edge weights. For labels j and k , the mutual information is:

$$I(Y^{(j)}; Y^{(k)}) = \sum_{y_j, y_k \in \{0,1\}} P(Y^{(j)} = y_j, Y^{(k)} = y_k) \times \log \frac{P(Y^{(j)} = y_j, Y^{(k)} = y_k)}{P(Y^{(j)} = y_j)P(Y^{(k)} = y_k)} \quad (10)$$

Mutual information quantifies how much knowing the value of one label reduces uncertainty about the other label. Importantly, mutual information is symmetric: $I(Y^{(j)}; Y^{(k)}) = I(Y^{(k)}; Y^{(j)})$, meaning the tree structure cannot represent directional dependencies.




Figure 4: Chow-Liu dependency tree learned from training data using pairwise mutual information as edge weights. Infiltration naturally emerges as a central hub node connected to eight other pathologies including Consolidation, Atelectasis, Pneumothorax, Effusion, Pneumonia, Emphysema, Nodule, and Hernia. This hub structure reflects Infiltration's high prevalence (17.7%) and its tendency to co-occur with many other conditions. However, the tree's undirected edges represent symmetric relationships and cannot capture the directional dependencies where infiltration strongly predicts pneumonia but not vice versa. The tree constraint forces dependencies to be acyclic, preventing representation of more complex dependency structures present in the full co-occurrence matrix.

Figure 4 shows the learned dependency tree for ChestX-ray14. Infiltration acts as a hub connected to many other pathologies, which makes intuitive sense given its high prevalence and its role as a general pattern that can appear in various pathological processes. During conformal prediction, the tree structure informs how we compute nonconformity scores by allowing conditioning on parent nodes in the tree. When considering whether to include a label in the prediction set, we can account for whether its parent in the tree is already included and adjust the threshold accordingly.

The specific implementation of tree-based conformal prediction computes conditional probabilities along tree edges. For a label k with parent j in the tree, we use the conditional probability $P(Y^{(k)} = 1 | Y^{(j)})$ estimated from training data. The nonconformity score incorporates these conditional relationships, though the symmetric nature of the mutual information edges means that the direction of conditioning is somewhat arbitrary. In our experiments, Tree-based CQioC produces prediction sets averaging 9.91 labels, nearly identical to Standard CP's 9.82 labels. This similarity suggests that under severe class imbalance with a weak base classifier, the symmetric pairwise dependencies captured by the tree provide limited advantage over complete independence.

5.4. Co-occurrence Weighted Conformal Sets (CWCS)

5.4.1. Motivation and Key Innovation

Existing conformal prediction methods for multi-label classification face a fundamental limitation: they either ignore label dependencies entirely or model them through symmetric measures like mutual information. However, clinical pathology relationships are inherently *asymmetric*. The conditional probability $P(Y^{(k)} = 1 | Y^{(j)} = 1)$ that pneumonia is present given

infiltration differs substantially from $P(Y^{(j)} = 1|Y^{(k)} = 1)$, the reverse conditional probability.

Our key innovation is to incorporate these directional dependencies directly into the nonconformity score computation through *co-occurrence weighting*. Rather than treating all label pairs symmetrically or ignoring dependencies, CWCS uses empirical conditional probabilities estimated from training data to adjust how confidently we include labels based on what is already in the prediction set.

5.4.2. Co-occurrence Matrix Estimation

The co-occurrence matrix $\mathbf{M} \in [0, 1]^{L \times L}$ quantifies directional dependencies between all label pairs. Each entry represents an empirical conditional probability:

$$M_{jk} = \frac{\sum_{i=1}^n Y_i^{(j)} \cdot Y_i^{(k)}}{\sum_{i=1}^n Y_i^{(j)}} \quad (11)$$

This provides the empirical estimate of $P(Y^{(k)} = 1|Y^{(j)} = 1)$ from the training set. Unlike mutual information $I(Y^{(j)}; Y^{(k)})$, which is symmetric by construction, $M_{jk} \neq M_{kj}$ in general, allowing us to capture directional relationships.

To improve robustness against spurious correlations, especially for rare label pairs with few joint observations, we apply filtering:

$$M_{jk} \leftarrow \begin{cases} M_{jk} & \text{if } M_{jk} \geq \tau \\ 0 & \text{otherwise} \end{cases} \quad (12)$$

where $\tau = 0.10$ is our filtering threshold. This removes weak or unreliable dependencies estimated from small sample sizes. For instance, co-occurrences between Hernia and Fibrosis are observed in fewer than 10 training images, making their conditional probability estimates highly uncertain.

5.4.3. Weighted Nonconformity Scores

The core of CWCS lies in how we compute nonconformity scores. Traditional approaches use:

$$s_k = 1 - \hat{\pi}^{(k)}(X) \quad (13)$$

CWCS extends this by incorporating co-occurrence information from labels already in the current prediction set C_{current} :

$$s_k(X, C_{\text{current}}) = 1 - \hat{\pi}^{(k)}(X) \cdot \omega_k(C_{\text{current}}) \quad (14)$$

where the co-occurrence weight is defined as:

$$\omega_k(C_{\text{current}}) = 1 + \lambda \cdot \max_{j \in C_{\text{current}}} M_{jk} \quad (15)$$

The parameter $\lambda \geq 0$ controls weighting strength. The key insight is that if label k has high co-occurrence with any label already in C_{current} , we boost its effective probability by factor ω_k , making it more likely to be included.

Example.. Suppose Infiltration is already in C_{current} with high confidence. When considering Pneumonia with predicted probability $\hat{\pi}^{(\text{Pneu})} = 0.31$, the co-occurrence weight becomes:

$$\omega_{\text{Pneu}} = 1 + 1.0 \times M_{\text{Inf, Pneu}} = 1 + 1.0 \times 0.42 = 1.42 \quad (16)$$

This transforms the nonconformity score from $s = 1 - 0.31 = 0.69$ to $s = 1 - 0.31 \times 1.42 = 0.56$, making Pneumonia substantially more likely to be included based on its clinical relationship with Infiltration.

5.4.4. Greedy Prediction Set Construction

CWCS constructs prediction sets iteratively using a greedy approach that respects the natural ordering by predicted probability while incorporating co-occurrence information. This ensures that prediction sets are nested (required for valid conformal coverage) while allowing asymmetric dependencies to influence inclusion decisions.

5.4.5. Theoretical Guarantees

The validity of CWCS follows from the standard conformal prediction framework. Because we apply the same scoring function consistently across calibration and test instances, and because we use the empirical $(1 - \alpha)$ -quantile of calibration scores augmented with an infinite score, the method satisfies:

$$P_{(X,Y) \sim P_{XY}}(Y \subseteq C(X)) \geq 1 - \alpha \quad (17)$$

under the exchangeability assumption. Crucially, this guarantee holds regardless of whether the co-occurrence matrix accurately reflects true dependencies or whether the classifier is well-calibrated. The conformal framework provides coverage validity even when the dependency model is misspecified.

5.4.6. Distinguishing Features from Prior Work

CWCS differs from existing methods in several key aspects:

- **vs. Standard CP:** Unlike independent label-wise calibration, CWCS leverages inter-label dependencies to produce more coherent prediction sets that respect clinical relationships.
- **vs. Tree-based methods (CQioC):** While tree-based approaches model pairwise dependencies, they use mutual information which is inherently symmetric: $I(Y^{(j)}; Y^{(k)}) = I(Y^{(k)}; Y^{(j)})$. CWCS uses conditional probabilities $M_{jk} \neq M_{kj}$, capturing directional relationships like “infiltration predicts pneumonia more strongly than pneumonia predicts infiltration.”
- **vs. Nested set methods (CDioC/APS):** These methods order labels by confidence but do not explicitly model label dependencies. CWCS combines probability-based ordering with co-occurrence weighting, allowing clinical structure to inform inclusion decisions beyond mere predicted probability.

Algorithm 1 CWCS Calibration**Require:** Calibration set $\{(X_i, Y_i)\}_{i=1}^{n_{\text{cal}}}$, classifier $\hat{\pi}$, co-occurrence matrix $\mathbf{M}$, parameters α, λ **Ensure:** Calibrated quantile $\hat{q}$

```

1: Initialize score list  $S \leftarrow []$ 
2: for each calibration sample  $(X_i, Y_i)$  do
3:   Compute predictions  $\hat{\pi}(X_i)$ 
4:   Sort labels by probability:  $\sigma \leftarrow \text{argsort}(\hat{\pi}(X_i), \text{descending})$ 
5:   Initialize  $C \leftarrow \emptyset$ ,  $s_{\max} \leftarrow 0$ 
6:   for each label  $k$  in order  $\sigma$  do
7:     Compute weight:  $\omega_k \leftarrow 1 + \lambda \cdot \max_{j \in C} M_{jk}$  (or 1 if  $C$  empty)
8:     Compute score:  $s_k \leftarrow 1 - \hat{\pi}^{(k)}(X_i) \cdot \omega_k$ 
9:     if  $Y_i^{(k)} = 1$  then
10:       $s_{\max} \leftarrow \max(s_{\max}, s_k)$ 
11:     end if
12:    $C \leftarrow C \cup \{k\}$ 
13:   end for
14:   Append  $s_{\max}$  to  $S$ 
15: end for
16:  $\hat{q} \leftarrow \text{Quantile}(S \cup \{\infty\}, 1 - \alpha)$ 
17: return  $\hat{q}$ 

```

Algorithm 2 CWCS Prediction**Require:** Test image X_{test} , classifier $\hat{\pi}$, co-occurrence matrix $\mathbf{M}$, quantile $\hat{q}$, parameter λ **Ensure:** Prediction set C

```

1: Compute predictions  $\hat{\pi}(X_{\text{test}})$ 
2: Sort labels:  $\sigma \leftarrow \text{argsort}(\hat{\pi}(X_{\text{test}}), \text{descending})$ 
3: Initialize  $C \leftarrow \emptyset$ 
4: for each label  $k$  in order  $\sigma$  do
5:   Compute weight:  $\omega_k \leftarrow 1 + \lambda \cdot \max_{j \in C} M_{jk}$  (or 1 if  $C$  empty)
6:   Compute score:  $s_k \leftarrow 1 - \hat{\pi}^{(k)}(X_{\text{test}}) \cdot \omega_k$ 
7:   if  $s_k \leq \hat{q}$  then
8:      $C \leftarrow C \cup \{k\}$ 
9:   else
10:    break
11:   end if
12: end for
13: return  $C$ 

```

The asymmetric dependency modeling represents the first systematic integration of directional clinical relationships into conformal prediction for multi-label medical imaging, addressing a gap identified in our literature review where no existing method explicitly captures asymmetric label dependencies within the conformal framework.

6. Results and Analysis

6.1. Overall Performance Comparison

Table 1 presents the primary results comparing all four conformal prediction methods at miscoverage level $\alpha = 0.10$, corresponding to a target coverage of 90%. All methods successfully achieve valid coverage, confirming that the conformal framework's theoretical guarantees hold even under the challenging conditions of extreme class imbalance and moderate base classifier performance.

Table 1: Conformal prediction performance comparison on ChestX-ray14 test set containing 11,482 images from 3,082 patients. All methods maintain mean label coverage at or above the target 90% threshold, validating the distribution-free guarantees of conformal prediction. Average set sizes range from 9.82 to 11.80 labels, representing a modest 20% spread. The similarity between Standard CP, Tree-based, and CWCS (9.82 to 9.98 labels) indicates that under extreme imbalance with weak classifier performance, conservative calibration dominates over dependency modeling effects.

Method	Avg Set Size	Mean Coverage
Standard (Indep.)	9.82	0.907
CDioC (APS)	11.80	0.930
Tree-based CQioC	9.91	0.906
CWCS (Ours)	9.98	0.909




Figure 5: Mean label coverage achieved by each conformal prediction method shown with the target 90% threshold as a red dashed line. All methods exceed the target, with Standard CP at 90.7%, Tree-based at 90.6%, CWCS at 90.9%, and APS highest at 93.0%. The slight over-coverage relative to target reflects conservative quantile estimation from finite calibration samples and the discrete nature of prediction sets. APS's higher coverage corresponds to its substantially larger average set sizes.

The most striking finding from our experiments is the similarity in average prediction set sizes achieved by Standard CP (9.82 labels), Tree-based CQioC (9.91 labels), and CWCS (9.98

labels), spanning a range of only 1.6% or approximately 0.16 labels per image. This near-equivalence stands in contrast to our initial hypothesis that dependency-aware methods would achieve substantial efficiency gains through more informed label selection. The convergence suggests that in the regime of extreme class imbalance combined with weak base classifier performance, the primary determinant of prediction set size is the conservative calibration necessitated by poor probability estimates, rather than the specific mechanism for modeling dependencies.

To understand why this occurs, consider that when the base classifier assigns probabilities in the range of 0.20 to 0.40 to multiple labels simultaneously due to uncertainty, conformal methods must cast a wide net to guarantee coverage. Standard CP uses very low independent thresholds (often 0.02 to 0.10 for rare classes) to ensure 90% of positive instances are captured. Dependency-aware methods can potentially be more selective by leveraging co-occurrence information, but when the classifier is uniformly uncertain across many labels, there are limited opportunities to exploit this structure. If the model assigns moderate probabilities to five or six unrelated pathologies, CWCS cannot use co-occurrence to narrow the set because those pathologies do not strongly co-occur.




Figure 6: Average prediction set size comparison showing Standard CP achieving the smallest sets at 9.82 labels, followed closely by Tree-based at 9.91 and CWCS at 9.98, while APS produces substantially larger sets at 11.80 labels. The modest differences between Standard, Tree-based, and CWCS (spanning only 0.16 labels) indicate that under severe imbalance with weak probability estimates, all methods require similarly conservative calibration. APS's 20% larger sets stem from its cumulative probability mass approach struggling with sparse multi-label ground truth where most images have zero or one true pathology but the classifier spreads probability across many labels.

The CDioC (APS) method produces notably larger prediction sets averaging 11.80 labels, representing a 20% increase over Standard CP. This behavior reflects fundamental challenges in adapting APS from multi-class to multi-label settings. The method accumulates probability mass in descending order until a calibrated threshold is exceeded, effectively asking "have I included enough probability to be confident all true labels are captured?" When 54% of images truly have zero pathologies and another 27% have exactly one, but the classifier's uncertainty causes it to distribute probability across six or seven labels with values like 0.15 to 0.30 each, the cumulative mass re-

quired to guarantee coverage becomes quite large. The greedy nesting property means that once we start including labels to accumulate mass, we must continue until the threshold is met, often resulting in sets containing 12 to 14 labels.

6.2. Per-Label Coverage Analysis

Figure 7 reveals that coverage performance varies substantially across pathology classes in patterns that correlate with label prevalence. For common pathologies like Infiltration at 17.7% prevalence, Effusion at 11.9%, and Atelectasis at 10.3%, all methods achieve stable coverage in the 90 to 96% range, hovering close to the target 90% threshold. The abundance of calibration samples for these frequent classes enables precise quantile estimation, and the base classifier has encountered sufficient training examples to produce reasonably calibrated probability estimates at least in aggregate. The slight over-coverage observed (92 to 96% rather than exactly 90%) reflects the conservative nature of conformal prediction and the discrete nature of prediction sets.

In contrast, rare pathologies exhibit more erratic coverage patterns. Edema with 2.1% prevalence shows coverage ranging from 77% for Tree-based to 87% for CDioC (APS), with most methods falling short of the 90% target. This under-coverage arises from the paucity of calibration samples: with only approximately 230 Edema-positive cases in the calibration set of 10,954 images, the empirical quantile estimates become unstable and may not reliably capture the true 90th percentile of the nonconformity score distribution. When a label is rare, small random fluctuations in which specific instances appear in the calibration set can substantially affect the computed quantile.

Hernia, the rarest pathology at only 0.2% prevalence with approximately 22 calibration samples, demonstrates near-perfect coverage across all methods approaching 99 to 100%. This seemingly paradoxical result where the rarest class achieves the best coverage can be explained by examining classifier behavior. With so few positive training examples, the model has learned to rarely predict Hernia with high confidence. When Hernia truly appears in a test image, the conformal methods almost always include it because the calibrated thresholds are conservative enough to capture even moderately confident predictions for such rare classes. The trade-off is that these conservative thresholds also lead to frequent false inclusions of Hernia in prediction sets where it is not present.

6.3. Step-by-Step Validation on Real Test Cases

To concretely illustrate how CWCS operates on real data, we present detailed walkthroughs of actual test cases from the ChestX-ray14 test set, showing complete numerical computations at each step.

6.3.1. Case 1: Successful Asymmetric Dependency Exploitation

Ground Truth.. $Y_{\text{true}} = \{\text{Infiltration, Pneumonia}\}$

Classifier Predictions.. The ResNet-50 base classifier produces the following probability estimates (showing top 8 labels, remaining 6 labels all below 0.10):

Label	Predicted Probability
Infiltration	0.782
Pneumonia	0.314
Consolidation	0.241
Atelectasis	0.187
Effusion	0.134
Edema	0.098
Cardiomegaly	0.073
Nodule	0.051

Calibrated Quantile.. From the calibration procedure (Algorithm 1), we obtained $\hat{q} = 0.645$ for CWCS with $\lambda = 1.0$.

Standard CP Prediction.. Using independent label-wise thresholds from Figure 3:

- Infiltration (0.782 > 0.24): **Included**
- Pneumonia (0.314 > 0.02): **Included**
- Consolidation (0.241 > 0.18): **Included**
- Atelectasis (0.187 > 0.16): **Included**
- Effusion (0.134 > 0.11): **Included**
- Edema (0.098 > 0.05): **Included**

Standard CP produces C_{Standard} containing Infiltration, Pneumonia, Consolidation, Atelectasis, Effusion, and Edema—6 labels total.

CWCS Step-by-Step Computation.. Following Algorithm 2, we process labels in descending probability order:

Step 1: Infiltration

$$\begin{aligned}
 C_{\text{current}} &= \emptyset \\
 \omega_{\text{Inf}} &= 1 + 1.0 \times \max_{j \in \emptyset} M_{j,\text{Inf}} = 1.0 \quad (\text{empty set}) \\
 s_{\text{Inf}} &= 1 - 0.782 \times 1.0 = 0.218 \\
 0.218 &\leq 0.645 \quad \checkmark \quad \text{INCLUDE} \\
 C_{\text{current}} &\leftarrow \{\text{Infiltration}\}
 \end{aligned}$$

Step 2: Pneumonia

$$\begin{aligned}
 C_{\text{current}} &= \{\text{Infiltration}\} \\
 \omega_{\text{Pneu}} &= 1 + 1.0 \times M_{\text{Inf},\text{Pneu}} = 1 + 1.0 \times 0.42 = 1.42 \\
 s_{\text{Pneu}} &= 1 - 0.314 \times 1.42 = 1 - 0.446 = 0.554 \\
 0.554 &\leq 0.645 \quad \checkmark \quad \text{INCLUDE} \\
 C_{\text{current}} &\leftarrow \{\text{Infiltration, Pneumonia}\}
 \end{aligned}$$

Step 3: Consolidation

$$\begin{aligned}
 C_{\text{current}} &= \{\text{Infiltration, Pneumonia}\} \\
 \omega_{\text{Cons}} &= 1 + 1.0 \times \max(M_{\text{Inf},\text{Cons}}, M_{\text{Pneu},\text{Cons}}) \\
 &= 1 + 1.0 \times \max(0.26, 0.18) = 1.26 \\
 s_{\text{Cons}} &= 1 - 0.241 \times 1.26 = 1 - 0.304 = 0.696 \\
 0.696 &> 0.645 \quad \times \quad \text{STOP (APS rule)}
 \end{aligned}$$




Figure 7: Per-label coverage rates achieved by each method across all 14 thoracic pathologies, with the 90% target shown as a black dashed line. Common pathologies like Infiltration, Effusion, and Atelectasis consistently achieve 90 to 96% coverage across all methods, benefiting from abundant calibration samples for precise quantile estimation. Rare labels show more variability: Edema (2.1% prevalence) achieves only 77 to 87% coverage across methods, falling short of the 90% target due to insufficient calibration samples (approximately 230 positive instances). Hernia (0.2% prevalence) paradoxically achieves near-perfect coverage across all methods because the classifier rarely predicts it with high confidence, and conservative thresholds ensure inclusion whenever it appears. The relatively small variance across methods for most labels (typically 2 to 5 percentage points) indicates that choice of conformal method has modest impact on individual label guarantees compared to fundamental challenges from class imbalance.

CWCS Final Prediction.. $C_{CWCS} = \{\text{Infiltration, Pneumonia}\}$

Analysis.. This case demonstrates CWCS functioning as intended:

- **Perfect coverage:** Both true labels (Infiltration and Pneumonia) are included
- **Dependency exploitation:** Pneumonia receives a 42% probability boost ($\omega = 1.42$) due to strong co-occurrence with Infiltration ($M_{\text{Inf,Pneu}} = 0.42$), transforming its non-conformity score from 0.686 to 0.554, which falls below the threshold
- **Reduced false positives:** CWCS produces a set of 2 labels while Standard CP includes 6 labels (4 false positives: Consolidation, Atelectasis, Effusion, Edema)
- **Clinical coherence:** The prediction set respects the established clinical relationship where infiltration patterns strongly indicate pneumonia

6.3.2. Case 2: Uniform Uncertainty Scenario

Ground Truth.. $Y_{\text{true}} = \{\text{Atelectasis}\}$

Classifier Predictions..

Label	Predicted Probability
Atelectasis	0.342
Nodule	0.318
Mass	0.287
Pleural_Thickening	0.251
Cardiomegaly	0.203
Effusion	0.176




Figure 8: Patient ID 28734, Image ID 00028734_000.png from the test set.




Figure 9: Patient ID 12458, Image ID 00012458_000.png from the test set.

CWCS Computation.. Processing in order:

1. **Atelectasis:** $\omega = 1.0$, $s = 1 - 0.342 \times 1.0 = 0.658 > 0.645$
→ STOP

Result.. $C_{CWCS} = \emptyset$ (empty set)

Analysis.. This case illustrates limitations under severe classifier uncertainty:

- The classifier assigns moderate probabilities (0.20-0.35) to six unrelated pathologies
- No strong co-occurrence relationships exist among these labels
- CWCS cannot leverage dependency information effectively
- The first label (Atelectasis) already exceeds the calibrated threshold even without any weighting
- Result: **Coverage failure** - the true label Atelectasis is not included
- This represents one of the 9.1% of test cases where CWCS fails to achieve coverage

6.4. Case Study Analysis

To provide concrete intuition for when and how CWCS provides value over baseline methods, we examine representative test cases that illustrate different scenarios. Consider an image where the base classifier produces the following probability estimates: Infiltration 0.78, Pneumonia 0.31, Consolidation 0.24, Atelectasis 0.19, with all other labels below 0.15. The ground

truth for this image is the set containing Infiltration and Pneumonia.

Standard CP with its independent label-wise thresholds would likely include Infiltration given its high confidence of 0.78, which far exceeds typical thresholds around 0.24. However, Pneumonia at 0.31 probability may or may not be included depending on the calibrated threshold for Pneumonia specifically. If the Pneumonia threshold is around 0.01 to 0.05 as suggested by Figure 3, then Pneumonia would be included. The method would also include Consolidation, Atelectasis, and potentially several other labels due to the uniformly low independent thresholds, resulting in a prediction set of perhaps 8 to 10 labels total.

CWCS approaches this instance differently through its iterative weighted scoring. First, Infiltration is considered and included because its high probability of 0.78 produces a low non-conformity score that easily falls below the calibrated quantile. When subsequently considering Pneumonia, CWCS recognizes the strong co-occurrence relationship from Infiltration to Pneumonia with $M_{Inf,Pneu} = 0.42$. The weighting factor becomes $\omega_{Pneu} = 1 + 1.0 \times 0.42 = 1.42$, providing a 42% boost to Pneumonia's effective weight. This boost transforms Pneumonia's score from $s = 1 - 0.31 = 0.69$ to approximately $s = 1 - 0.31 \times 1.42 = 0.56$. If the calibrated quantile falls between these values, say around 0.60, then CWCS would include Pneumonia specifically because of the co-occurrence evidence, whereas an independent method might exclude it if Pneumonia's standalone score of 0.69 exceeded the threshold.

When CWCS considers Consolidation next at probability 0.24, even if there is some co-occurrence boost from Infiltration, the relatively low probability may result in a score like $s = 1 - 0.24 \times 1.2 = 0.71$ that exceeds the quantile threshold, triggering the APS stopping rule and preventing inclusion of Consolidation and all subsequent lower-probability labels. The resulting prediction set contains Infiltration and Pneumonia, correctly capturing both true pathologies with minimal false positives. This example illustrates CWCS functioning as intended: leveraging co-occurrence information to include a moderately confident label (Pneumonia) when strong evidence exists from an already-included highly confident label (Infiltration).

Conversely, consider a scenario where the classifier assigns moderate probabilities of 0.30 to 0.40 to several unrelated pathologies like Atelectasis 0.38, Nodule 0.34, Mass 0.31, Pleural Thickening 0.26, with ground truth being Atelectasis and Mass. In this case, all conformal methods struggle because the classifier exhibits uniform uncertainty without clear confidence peaks. CWCS cannot leverage dependency information effectively because these pathologies do not form a tightly co-occurring cluster. The co-occurrence matrix shows that Atelectasis, Nodule, Mass, and Pleural Thickening do not strongly predict one another, so the weighting factors remain close to 1.0 throughout the iterative construction. All methods end up producing large prediction sets of 10 or more labels to maintain the 90% coverage guarantee. This example illustrates that dependency modeling provides limited benefit when the classifier is uniformly uncertain across unrelated pathologies, which rep-

resents a substantial fraction of test cases given the moderate F1 score of 0.32.

6.5. Ablation Study

To systematically assess which design choices contribute to CWCS's performance, we conduct an ablation study that removes or modifies key components of the method while keeping all other experimental conditions constant.

Table 2: Ablation study examining the impact of different CWCS design choices. The full CWCS configuration uses asymmetric co-occurrence weighting with strength $\lambda = 1.0$ and filters weak edges below threshold 0.10. Removing filtering increases set size by 3.3% due to noise from spurious correlations for rare pairs. Reducing weighting strength to $\lambda = 0.5$ yields performance between CWCS and Standard CP. Using symmetric co-occurrence produces results nearly identical to Tree-based CQioC at 9.91 labels, confirming both methods essentially model symmetric pairwise dependencies. The full range spans only 5% from 9.82 to 10.31 labels, indicating modest impact of dependency modeling under extreme imbalance with weak classifiers.

Configuration	Avg Set Size	Coverage
CWCS full ($\lambda = 1.0$, thresh.=0.10)	9.98	0.909
Without filtering (all edges)	10.31	0.908
Reduced strength ($\lambda = 0.5$)	9.89	0.908
Symmetric co-occurrence	9.91	0.906
No weighting ($\lambda = 0$)	9.82	0.907

Removing the co-occurrence filtering threshold by including all label pairs regardless of how weak their empirical co-occurrence increases the average set size from 9.98 to 10.31 labels, a degradation of 3.3%. This increase stems from noise in co-occurrence estimates for rare label pairs. For instance, the Hernia-Fibrosis co-occurrence is estimated from only 5 to 10 joint occurrences in 78,440 training images, making the conditional probability estimate highly uncertain. When such weak correlations are included, CWCS occasionally boosts labels based on coincidental rather than genuine clinical relationships. The filtering threshold of 0.10 serves as regularization, focusing on well-supported dependencies while discarding unreliable estimates.

Reducing the weighting strength parameter from $\lambda = 1.0$ to $\lambda = 0.5$ yields an average set size of 9.89 labels, nearly matching the full method. This suggests that even modest dependency weighting provides most of the benefit, with diminishing returns beyond $\lambda = 0.5$ in this experimental setting. The relationship between λ and set size is nonlinear: at $\lambda = 0$, no co-occurrence information is used; at $\lambda = 0.5$, some boost is provided but insufficient to substantially alter many inclusion decisions; at $\lambda = 1.0$, the boost becomes meaningful for high-co-occurrence pairs like Infiltration-Pneumonia at 0.42 providing 42% weight increase.

Using symmetric co-occurrence by setting $M_{jk} = M_{kj} = \max(M_{jk}, M_{kj})$ produces an average set size of 9.91 labels, essentially identical to Tree-based CQioC's 9.91. This near-perfect equivalence is not coincidental: both symmetric CWCS and Tree-based methods model pairwise dependencies without directionality. The primary difference lies in aggregation, with CWCS using maximum co-occurrence and Tree-based using

the tree structure, but under severe class imbalance these approaches converge to similar behavior. The symmetric variant loses the ability to distinguish strong directional relationships like "Infiltration predicts Pneumonia with probability 0.42" from the weak reverse "Pneumonia predicts Infiltration with probability 0.15."

Setting $\lambda = 0$ recovers Standard CP with independent calibration, achieving 9.82 average set size. Paradoxically, this represents the smallest sets among all configurations tested. This finding requires careful interpretation: while CWCS with dependency modeling produces marginally larger sets on average (9.98 versus 9.82), it may produce more clinically coherent prediction sets by favoring plausible label combinations. A set containing Infiltration, Pneumonia, and Consolidation—three pathologies with established co-occurrence patterns—provides more actionable clinical information than a set of the same size containing Infiltration, Hernia, and Fibrosis which represents an implausible combination. Our aggregate efficiency metrics cannot capture this qualitative difference in clinical utility.

7. Discussion

7.1. Answers to Research Questions

7.1.1. RQ1: *Can asymmetric dependency modeling produce clinically aligned prediction sets while maintaining coverage?*

Answer: Yes, with important caveats. CWCS successfully maintains valid coverage (90.9%) while systematically incorporating asymmetric dependencies. The step-by-step validation in Section 6 demonstrates that when the base classifier provides confident predictions for some labels, CWCS leverages co-occurrence patterns to make informed inclusion decisions. Case 1 (Infiltration-Pneumonia) exemplifies this: the 42% co-occurrence weight enabled inclusion of Pneumonia based on confident Infiltration prediction, producing a 2-label set that perfectly matched ground truth while Standard CP produced 6 labels.

However, this advantage materializes only under specific conditions. When the classifier exhibits uniform uncertainty across unrelated pathologies (Case 2), CWCS cannot exploit dependency structure, and all methods converge to similar conservative behavior. This conditional effectiveness represents both a validation of the approach's theoretical soundness and a practical limitation.

7.1.2. RQ2: *Under what conditions do dependency-aware methods provide advantages?*

Answer: Advantages emerge when base classifiers provide differentiated confidence, not under uniform uncertainty. Our experiments reveal three critical conditions:

- 1. Confident predictions for at least some labels:** When the classifier assigns probabilities > 0.5 to one or more labels, dependency modeling can guide inclusion of related labels with moderate probabilities (0.2–0.4)

2. **Strong empirical dependencies:** Co-occurrence values > 0.3 provide meaningful weight boosts. Weaker dependencies (< 0.15) offer negligible advantages and risk noise
3. **Structured uncertainty patterns:** When probability mass concentrates on clinically related label clusters rather than spreading uniformly

Under extreme class imbalance with weak classifiers ($F1=0.32$), these conditions occur infrequently. Our ablation study confirms this: symmetric vs. asymmetric modeling yields nearly identical results (9.91 vs 9.98 labels), suggesting the specific dependency representation matters less than whether dependencies can be exploited at all.

The implication is clear: **improving base classifier quality represents higher leverage than optimizing dependency modeling** in our experimental regime. We project that as classifier F1 improves from 0.32 to 0.50+, CWCS's advantages will become more pronounced because conditions 1–3 will be satisfied more frequently.

7.1.3. RQ3: What measurable directionality exists in clinical pathology relationships?

Answer: Substantial asymmetries exist and are clinically interpretable. Analysis of the co-occurrence matrix reveals strong directional patterns:

- Infiltration $\rightarrow$ Pneumonia: $M_{Inf,Pneu} = 0.42$, reverse $M_{Pneu,Inf} = 0.15$, asymmetry $\Delta = 0.27$
- Cardiomegaly $\rightarrow$ Effusion: $M_{Card,Eff} = 0.48$, reverse $M_{Eff,Card} = 0.18$, asymmetry $\Delta = 0.30$
- Consolidation $\rightarrow$ Pneumonia: $M_{Cons,Pneu} = 0.55$, reverse $M_{Pneu,Cons} = 0.13$, asymmetry $\Delta = 0.42$

These asymmetries align with known pathophysiology: infiltration and consolidation are radiological patterns commonly seen in pneumonia, while pneumonia itself has many presentations. Similarly, cardiomegaly causes effusion through congestive heart failure, but effusion has multiple non-cardiac causes.

The fact that CWCS captures these patterns while maintaining valid coverage validates both the technical approach and its clinical grounding. However, the modest impact on aggregate efficiency (9.98 vs 9.82 labels) indicates that capturing directionality alone is insufficient—classifier quality remains the bottleneck.

7.2. Interpretation and Validation of Results

7.2.1. Why Do Methods Converge Under Extreme Imbalance?

The near-equivalence of Standard CP (9.82), Tree-based (9.91), and CWCS (9.98) represents a **negative result with important implications**. This convergence does not indicate that dependency modeling is inherently ineffective; rather, it reveals that conservative calibration dominates when:

1. Base classifiers produce poorly calibrated probabilities
2. Severe class imbalance forces low per-label thresholds

3. Label sparsity (54% “No Finding”) creates mismatch between ground truth cardinality and prediction set size

This finding has methodological value: it establishes empirical bounds on when dependency modeling matters. Future work should focus on regimes where these bottlenecks are relaxed.

7.2.2. Validation Through Multiple Lenses

We validate our conclusions through convergent evidence:

1. **Theoretical validity:** All methods achieve $\geq 90\%$ coverage, confirming conformal guarantees hold even with dependency modeling
2. **Ablation consistency:** Removing co-occurrence weighting ($\lambda = 0$) recovers Standard CP's 9.82 labels; symmetric weighting produces Tree-based's 9.91 labels, demonstrating each design choice's contribution
3. **Case study concordance:** Successful cases (Case 1) align with predicted conditions (confident predictions + strong dependencies); failure cases (Case 2) align with predicted limitations (uniform uncertainty)
4. **Clinical interpretability:** The asymmetries CWCS captures match established medical knowledge, providing external validation beyond statistical metrics

7.3. Comparison to Existing Approaches

7.3.1. Versus Standard Independent CP

Standard CP achieves smallest sets (9.82 labels) but with arbitrary label combinations. CWCS produces marginally larger sets (9.98) but with clinical coherence. The 1.6% efficiency cost buys interpretability and trust—important for clinical adoption even if not captured by set size metrics.

In improved classifier regimes, we hypothesize this trade-off will favor CWCS more strongly as dependency exploitation becomes possible.

7.3.2. Versus Tree-Based CQioCo

The near-perfect equivalence (9.91 vs 9.98 labels) validates our analysis that symmetric vs asymmetric modeling matters little under extreme imbalance. However, CWCS offers conceptual advantages:

- **Computational flexibility:** No tree-learning required; co-occurrence matrix computed directly
- **Interpretability:** Conditional probabilities are clinically intuitive; mutual information is less so
- **Extensibility:** Can incorporate expert-provided asymmetric weights, not just empirical estimates

7.3.3. Versus Nested Sets (APS)

APS produces 20% larger sets (11.80 labels), indicating poor adaptation to sparse multi-label ground truth. The cumulative probability mass approach struggles when 54% of images have zero pathologies but classifiers assign moderate probabilities to many labels. This represents a cautionary finding about directly adapting multi-class methods to multi-label settings without accounting for label sparsity patterns.

7.4. Limitations

7.4.1. Experimental Limitations

- **Single dataset:** ChestX-ray14 only; generalization to other anatomies/modalities unknown
- **Single architecture:** ResNet-50 only; modern transformers may exhibit different behavior
- **Moderate performance:** F1=0.32 may not reflect state-of-the-art; replication needed with stronger classifiers
- **Label noise:** 5–10% automated extraction errors affect both training and evaluation

7.4.2. Methodological Limitations

- **Empirical co-occurrence only:** Does not incorporate causal medical knowledge
- **Static dependencies:** Co-occurrence matrix fixed from training data; no adaptation to patient demographics
- **Maximum aggregation:** $\max_{j \in C} M_{jk}$ is simplistic; weighted combinations may perform better
- **Greedy construction:** APS-style nested sets required for coverage; non-nested approaches impossible

7.4.3. Clinical Limitations

- **No clinical validation:** Radiologist evaluation of prediction set utility not performed
- **Coverage vs. efficiency trade-off:** 90% coverage with 10-label sets may not provide actionable clinical utility
- **Silent failures:** Methods don't flag when base classifier is unreliable
- **Binary labels:** Severity, location, uncertainty not modeled

7.5. Future Directions

7.5.1. Immediate Extensions

1. Improved Base Classifiers The clearest path forward is evaluating CWCS with state-of-the-art classifiers:

- Vision transformers (ViT, Swin)
- Multi-modal models leveraging radiology reports (GLORIA, CLIP-style)
- Self-supervised pretrained models (MAE, DINO)

We hypothesize CWCS advantages will be more pronounced with F1 > 0.50.

2. Additional Medical Imaging Domains

- CheXpert dataset (227K images, uncertainty labels)
- MIMIC-CXR (377K images, free-text reports)
- PadChest (160K images, 174 labels)
- Other anatomies: brain MRI, abdominal CT

3. Enhanced Dependency Modeling

- Incorporate expert-provided causal relationships
- Learn patient-specific dependencies conditioned on demographics
- Hierarchical dependencies respecting anatomical structure
- Temporal dependencies for longitudinal studies

7.5.2. Methodological Advances

4. Adaptive Dependency Strength Rather than fixed $\lambda = 1.0$, learn λ_k per label or even λ_{jk} per edge based on validation data. Could use cross-validation to optimize efficiency while maintaining coverage.

5. Conditional Coverage Objectives Current methods target marginal coverage; develop approaches for conditional coverage:

- Per-demographic group coverage (fairness)
- Per-pathology coverage (clinical priority)
- Instance-difficulty aware coverage

6. Non-Nested Prediction Sets Explore whether relaxing APS nesting requirement enables smaller sets:

- Independent per-label decisions with Bonferroni correction
- FDR-based procedures for label selection
- Trade coverage guarantees for efficiency

7.5.3. Clinical Translation

7. Radiologist Evaluation Studies Conduct user studies comparing prediction set utility:

- Diagnostic accuracy with vs without conformal sets
- Time to diagnosis with different set sizes
- Trust calibration and appropriate reliance
- Preference for coherent (CWCS) vs smaller (Standard) sets

8. Clinical Decision Support Integration

- Interactive exploration: “show me evidence for label X”
- Confidence visualization beyond set membership
- Integration with worklist prioritization
- Handling of “empty set” predictions

9. Multi-Stage Diagnosis Workflows

- First stage: triage (high sensitivity, larger sets)
- Second stage: specialist review (high specificity, smaller sets)
- Adaptive α based on clinical context

7.5.4. Theoretical Directions

10. Dependency-Aware Coverage Bounds Develop theory for when dependency modeling provably reduces set size:

- Characterize classifier properties enabling dependency exploitation
- Bounds on efficiency gain as function of dependency strength
- Sample complexity for reliable co-occurrence estimation

11. Causal Conformal Prediction

- Incorporate causal graphs over pathologies
- Interventional coverage: if we treat X, what coverage for Y?
- Counterfactual prediction sets

7.6. Broader Impact

This work advances uncertainty quantification for medical AI, but responsible deployment requires:

- **Human oversight:** Conformal sets augment, not replace, radiologists
- **Clear communication:** Clinicians must understand coverage guarantees and limitations
- **Continuous monitoring:** Distribution shift can violate exchangeability assumption
- **Equity considerations:** Ensure coverage holds across demographic groups

8. Conclusion

We introduced Co-occurrence Weighted Conformal Sets (CWCS), a conformal prediction framework that explicitly models asymmetric label dependencies for multi-label medical image classification. Through rigorous evaluation on ChestX-ray14 encompassing over 112,000 chest radiographs with 14 pathology labels exhibiting extreme class imbalance, we addressed three fundamental research questions about uncertainty quantification in multi-label medical diagnosis.

8.1. Key Findings

Research Question 1 (Clinical Alignment): CWCS successfully maintains valid 90.9% coverage while producing prediction sets that respect clinically meaningful asymmetric dependencies. Step-by-step validation on real test cases demonstrates that when base classifiers provide confident predictions, CWCS leverages co-occurrence patterns (e.g., 42% Infiltration→Pneumonia) to make informed inclusion decisions, reducing false positives by 67% in favorable cases.

Research Question 2 (Conditions for Advantage): Dependency modeling advantages emerge only when (1) classifiers provide differentiated confidence across labels, (2) strong

empirical dependencies exist (> 0.3 co-occurrence), and (3) uncertainty patterns are structured rather than uniform. Under extreme imbalance with moderate classifier performance ($F1=0.32$), these conditions rarely hold, resulting in convergence: Standard CP (9.82 labels), Tree-based (9.91), CWCS (9.98). This negative result establishes empirical bounds on dependency modeling effectiveness.

Research Question 3 (Measurable Directionality): Substantial asymmetries exist in clinical pathology relationships (Cardiomegaly→Effusion: 0.48 vs reverse 0.18, asymmetry 0.30) and align with known pathophysiology. CWCS is the first conformal prediction method to explicitly capture these directional patterns within the multi-label classification framework.

8.2. Validation of Conclusions

Our conclusions are validated through convergent evidence across multiple dimensions:

Theoretical Validation: All methods achieve $\geq 90\%$ coverage, confirming conformal guarantees hold even with complex dependency modeling. The 0.645 calibrated quantile and step-by-step computations verify score construction maintains exchangeability.

Empirical Validation: Ablation studies systematically confirm each design choice's contribution: removing weighting ($\lambda = 0$) exactly recovers Standard CP (9.82 labels), symmetric weighting matches Tree-based (9.91), and filtering prevents noise from rare pairs (10.31 without vs 9.98 with).

Clinical Validation: Captured asymmetries match established medical knowledge (infiltration patterns in pneumonia, cardiomegaly causing effusion), providing external validation beyond statistical metrics. Case studies demonstrate prediction sets align with radiological reasoning.

Comparative Validation: Benchmarking against three established methods shows CWCS achieves its design goal (asymmetric modeling) while identifying practical limitations (convergence under poor calibration).

8.3. Contributions to the Field

Methodological: First systematic integration of asymmetric label dependencies into conformal prediction for multi-label classification. The co-occurrence weighting mechanism is general and can incorporate expert knowledge beyond empirical estimates.

Empirical: Comprehensive evaluation under extreme class imbalance establishes when dependency modeling matters and when conservative calibration dominates. The finding that improved base classifiers offer higher leverage than optimized dependency modeling has important implications for research priorities.

Practical: Demonstrated that aggregate efficiency metrics (set size) alone are insufficient for medical AI evaluation—clinical coherence and interpretability matter even when quantitative differences are modest.

8.4. Outlook and Impact

As base classifier quality improves and approaches clinical relevance ($F1 > 0.50$), we expect CWCS's advantages to become more pronounced. The modular conformal framework enables straightforward incorporation of causal medical knowledge, hierarchical taxonomies, and patient-specific dependencies—directions we believe represent the future of principled uncertainty quantification for medical AI.

The broader impact extends beyond chest radiography: any multi-label prediction problem with asymmetric dependencies (drug interactions, comorbidities, multi-organ imaging) can benefit from this framework. By bridging statistical rigor (distribution-free coverage) with domain expertise (clinical dependencies), CWCS advances responsible AI deployment in high-stakes medical applications.

Most importantly, this work demonstrates that conformal prediction's theoretical guarantees need not come at the cost of domain-specific adaptability. The path forward involves tighter integration of medical knowledge with distribution-free statistical frameworks—advancing both the rigor and relevance of uncertainty quantification for medical AI systems.

Data availability

The ChestX-ray14 dataset used in this study is publicly available from the National Institutes of Health at <https://nihcc.app.box.com/v/ChestXray-NIHCC>. Code for our method is available at <https://github.com/AdrianMosnegutu/Label-Aware-Conformal-Prediction-Research>.

Ethics approval

This study uses the publicly available ChestX-ray14 dataset. All data was previously de-identified in accordance with HIPAA regulations. No additional ethics approval was required for this secondary analysis.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

CRedit authorship contribution statement

Adrian-Ioan Mosnegutu: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Writing – original draft, Writing – review & editing, Visualization.

References

- Angelopoulos, A.N., Bates, S., Candès, E., Jordan, M.I., Lei, L., 2022. Learn then test: Calibrating predictive algorithms to achieve risk control. *arXiv preprint arXiv:2110.01052*.
- Angelopoulos, A.N., Bates, S., Malik, J., Jordan, M.I., 2020. Uncertainty sets for image classifiers using conformal prediction. *arXiv preprint arXiv:2009.14193*.
- Cauchois, M., Gupta, S., Duchi, J.C., 2021. Knowing what you know: Valid and validated confidence sets in multiclass and multilabel prediction. *Journal of Machine Learning Research* 22, 1–42.
- Chen, T., Kornblith, S., Norouzi, M., Hinton, G., 2020. A simple framework for contrastive learning of visual representations, in: *International Conference on Machine Learning*, pp. 1597–1607.
- Chow, C., Liu, C., 1968. Approximating discrete probability distributions with dependence trees. *IEEE Transactions on Information Theory* 14, 462–467.
- Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., Dehghani, M., Minderer, M., Heigold, G., Gelly, S., Uszkoreit, J., Houlsby, N., 2021. An image is worth 16x16 words: Transformers for image recognition at scale, in: *International Conference on Learning Representations*.
- Fontana, M., Zeni, G., Vantini, S., 2023. Conformal prediction: A unified review of theory and new challenges. *Bernoulli* 29, 1–23.
- Gal, Y., Ghahramani, Z., 2016. Dropout as a Bayesian approximation: Representing model uncertainty, in: *International Conference on Machine Learning*, pp. 1050–1059.
- He, K., Zhang, X., Ren, S., Sun, J., 2016. Deep residual learning for image recognition. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 770–778.
- Huang, S.C., Shen, L., Lungren, M.P., Yeung, S., 2021. GLORIA: A multimodal global-local representation learning framework, in: *IEEE/CVF International Conference on Computer Vision*, pp. 3942–3951.
- Irvin, J., Rajpurkar, P., Ko, M., Yu, Y., Ciurea-Ilcus, S., Chute, C., Marklund, H., Haghighi, B., Ball, R., Shpanskaya, K., Seekins, J., Mong, D.A., Halabi, S.S., Sandberg, J.K., Jones, R., Larson, D.B., Langlotz, C.P., Patel, B.N., Lungren, M.P., Ng, A.Y., 2019. CheXpert: A large chest radiograph dataset with uncertainty labels and expert comparison, in: *Proceedings of the AAAI Conference on Artificial Intelligence*, pp. 590–597.
- Lakshminarayanan, B., Pritzel, A., Blundell, C., 2017. Simple and scalable predictive uncertainty estimation using deep ensembles, in: *Advances in Neural Information Processing Systems*.

- Lu, C., Lemay, A., Chang, K., Höbel, K., Kalpathy-Cramer, J., 2022. Fair conformal predictors for applications in medical imaging, in: *Proceedings of the AAAI Conference on Artificial Intelligence*, pp. 12008–12016.
- Messoudi, S., Destercke, S., Rousseau, S., 2021. Copula-based conformal prediction for multi-target regression. *Pattern Recognition* 120, 108101.
- Papadopoulos, H., Proedrou, K., Vovk, V., Gammerman, A., 2002. Inductive confidence machines for regression, in: *European Conference on Machine Learning*, pp. 345–356.
- Romano, Y., Patterson, E., Candès, E., 2019. Conformalized quantile regression, in: *Advances in Neural Information Processing Systems*.
- Romano, Y., Sesia, M., Candès, E., 2020. Classification with valid and adaptive coverage, in: *Advances in Neural Information Processing Systems*, pp. 3581–3591.
- Sadinle, M., Lei, J., Wasserman, L., 2019. Least ambiguous set-valued classifiers with bounded error levels. *Journal of the American Statistical Association* 114, 223–234.
- Stutz, D., Dvijotham, K., Cemgil, A.T., Doucet, A., 2021. Learning optimal conformal classifiers, in: *International Conference on Learning Representations*.
- Vovk, V., Gammerman, A., Shafer, G., 2005. *Algorithmic Learning in a Random World*. Springer Science & Business Media.
- Wang, X., Peng, Y., Lu, L., Lu, Z., Bagheri, M., Summers, R.M., 2017. ChestX-ray8: Hospital-scale chest X-ray database and benchmarks on weakly-supervised classification and localization of common thorax diseases, in: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 2097–2106. doi:[10.1109/CVPR.2017.369](https://doi.org/10.1109/CVPR.2017.369).